



a

—●— Claude 3.5 Haiku —■— Gemini 3.5 Flash-Lite —▲— GPT-4o —◆— Mistral Large

b

Panel a scores all four predictors on identical rows, so the comparison is about how much history the behaviour uses, not about sample size; the ringed marker is each model's selected depth, and all four select memory-two. Panel b covers only the games a single memory-one rule cannot explain but two consecutive ones can (Claude 374, Gemini 329, GPT-4o 377, Mistral 273). The mode is round 2-3 for every model, with a second Gemini mode at round 9.